

Exploring Customer Feedback: A Study of Reviews at Pizza Hut, Sri Lanka*

Chenyiteng Han

April 18, 2024

This paper uses statistical analysis to explore customer reviews from a Pizza Hut in Sri Lanka, focusing on the relationship between star ratings and the likelihood and length of text feedback. We discovered that lower star ratings often come with more detailed textual reviews. This points to the idea that reviews, especially negative ones, hold insights for businesses looking to improve customer satisfaction and service quality. Through this study, we highlight patterns in consumer behavior, emphasizing how a deep dive into feedback can play a role in bettering customer experiences.

Table of contents

1	Introduction	1
2	Data	2
2.1	Raw Data	3
2.2	Cleaned Data	4
2.3	Data Summary	5
2.3.1	Frequency of Text Reviews Across Star Categories	5
2.3.2	Overview of Text Length in Reviews by Star Rating	6
2.4	Measurement	7
3	Models	8
3.1	Logistic Regression Model for Text Indicator and Star Ratings	8
3.2	The Multilevel Model for Text Length and Star Ratings	9
4	Results	11
4.1	The Relationship Between Text reviews and Star Ratings	11

*Code and data from this analysis are available at: <https://github.com/Hhnxxxxxx/Pizza-Hut-Reviews.git>.

4.2	The Different Length of Text Reviews Between Different Stars	13
5	Discussion	14
5.1	Overview of Works	14
5.2	Insights into Societal Behaviors	15
5.3	Reflecting on the Limitations	16
5.4	Future Directions and Unexplored Avenues	16
	Appendix	17
A	Model Diagnostics	17
A.1	Diagnostics for the Logistic Regression Model	17
A.2	Diagnostics for the Multilevel Model	17
	References	18

1 Introduction

In this paper, we explore the world of customer feedback, focusing on identifying trends within star ratings and text reviews for a popular Pizza Hut location in Sri Lanka. Given the vast amount of ratings and reviews that businesses receive every day, we wanted to dig deeper into what customers are really trying to say through their feedback. We looked at two main things: whether customers leave text reviews along with their star ratings, and if so, how long these reviews tend to be.

The estimand of our study is the effect of customer star ratings on the probability of leaving textual feedback and the length of such feedback at a Pizza Hut outlet in Sri Lanka. By using a logistic regression model and a multilevel negative binomial model, as discussed in Section 3 Models, we tried to understand how likely customers are to leave text reviews and how detailed they are, depending on their level of satisfaction. The findings, which we talk about in Section 4 Results, showed a strong inverse relationship between star ratings and both the likelihood and length of text reviews. Basically, the lower the star rating, the more probable it was for a customer to leave a review, and the longer that review was likely to be. This insight is really important because it highlights how customers who are less satisfied are more likely to give detailed feedback, which could be very useful for businesses looking to make improvements.

The lessons from this study could be useful for the retail and service sectors, as they show the importance of paying close attention to customer feedback. The rest of the paper, from Section 2 Data to Section 5 Discussion, goes over the data we used, both before and after processing, the methods of analysis we applied, the insights we got from the results, and our thoughts on the limitations and what we might want to look at in future research, like maybe using natural language processing to get even more out of the text reviews. By getting a

better grasp of these findings and what they say about customer behavior, businesses can do a better job of adjusting their services to match what their customers want and expect, leading to happier customers all around.

2 Data

The foundation of our analysis on Pizza Hut customer feedback in Sri Lanka was a dataset obtained from Kaggle, accessed through the Kaggle API (Kaggle (2022)) which allows for automated and direct data retrieval. In our R (R Core Team (2022)) environment, the data-verse package (Leeper (2022)) was instrumental in organizing and managing the dataset. Data preparation, including cleaning and transformation, was handled adeptly with the tidyverse collection of packages (Wickham and al. (2019)), particularly utilizing dplyr (Wickham and al. (2022)) for its powerful data manipulation capabilities.

Visualization and interpretation of the data were greatly enhanced by ggplot2 (Wickham (2022)) for creating informative graphics, while RColorBrewer (Neuwirth (2022)) provided aesthetic color schemes to enrich these visual representations. The knitr package (Xie (2014)) facilitated the seamless integration of R code within our dynamic report, allowing for an interactive presentation of results.

For the advanced statistical modeling, we employed rstanarm (Team (2022)), which provided a suite of tools for Bayesian inference in generalized linear models. This was complemented by tidybayes (Kay (2022b)), which enabled a tidy workflow when working with Bayesian models. The arrow package (Foundation (2022)) was pivotal in efficient data reading and writing, ensuring quick access to our dataset.

The modelsummary package (Arel-Bundock (2022)) allowed us to create detailed statistical tables, effectively summarizing the results of our complex models. To visualize distributions and express uncertainty in our data, ggdist (Kay (2022a)) proved invaluable. For handling the outputs from mixed models, broom.mixed (Bolker et al. (2022)) was utilized, allowing for the tidying and interpretation of mixed effects models. Collectively, these R packages provided a robust framework for our research, ensuring a high standard of accuracy and reproducibility.

2.1 Raw Data

This study utilizes the “Pizza Hut Reviews - Insights from One of Sri Lanka’s Pioneer Branches” dataset, extracted from Kaggle. The dataset is a treasure trove of customer sentiments, encompassing 4000 entries of evaluative feedback for the Pizza Hut branch located at Union Place, Colombo, which is noted for being among the first to establish Pizza Hut’s presence in Sri Lanka.

As depicted in Table 1, the dataset provides a snapshot into the customers’ appraisals and narratives. Each record consists of three columns that collectively offer a picture of consumer opinion.

- title: This column consistently identifies the venue as “Pizza Hut - Union Place,” situating the reviews within the specific context of this branch’s performance.
- stars: Reflecting the customer’s satisfaction, this field quantifies their experience on a 1 to 5 scale, with 1 indicating dissatisfaction and 5 signifying an excellent service encounter.
- text: Perhaps the most telling of the columns, this contains the actual text of the customer’s review. It ranges from concise accolades to detailed narratives, providing qualitative depth beyond the numerical rating.

While alternative datasets from different branches or regions could have been considered, the choice to focus on this specific dataset was intentional. The Union Place branch’s historic significance within the Sri Lankan context presents a unique opportunity to understand customer feedback dynamics where brand establishment is deep-rooted.

In presenting the raw data in Table 1, readers are invited to witness the unfiltered voice of the customer, which will later be dissected and analyzed to glean actionable insights. This approach ensures that data points are scrutinized and contributes to careful analysis that respects customer feedback.

Table 1: Sample of the Raw Dataset

Title	Stars	Text
Pizza Hut - Union Place	5	
Pizza Hut - Union Place	5	
Pizza Hut - Union Place	4	
Pizza Hut - Union Place	5	
Pizza Hut - Union Place	4	

Table 1: Sample of the Raw Dataset

Title	Stars	Text
Pizza Hut	5	
- Union		
Place		
Pizza Hut	5	Good place to visit to have some food
- Union		
Place		
Pizza Hut	5	Had a personal pan chicken tandoori, it was quite on point just like the
- Union		pizza hut one looks forward to. The staff was courteous and had parking
Place		space too.
Pizza Hut	5	
- Union		
Place		
Pizza Hut	4	
- Union		
Place		

2.2 Cleaned Data

Following the data curation process, we have derived a cleaner and more analytically suitable dataset titled “clean_pizza_hut_reviews,” specifically sculpted to facilitate quantitative analysis. The original verbose feedback has been distilled into quantitative variables that retain the essence of customer ratings while allowing for numerical manipulation.

As demonstrated in Table 2, the refined dataset retains the granularity of ratings through the ‘stars’ variable, indicative of customer satisfaction levels, while introducing two new variables: ‘text_indicator’ and ‘words.’

- review_id: Assigned as a unique identifier to each review, this variable serves a structural purpose, facilitating reference and management within the dataset rather than bearing any intrinsic meaning.
- stars: Preserved from the raw data, this metric remains unchanged, symbolizing the customer’s rating, ranging from 1 for dissatisfaction to 5 for exemplary service.
- text_indicator: A binary variable, it denotes the presence (1) or absence (0) of textual commentary accompanying the star rating. This distinction is useful for subsequent analyses that differentiate between ratings supplemented with text and those that are not.

- words: This numeric variable quantifies the length of the text accompanying the review. It operates in tandem with the ‘text_indicator’—if a review is textually barren (‘text_indicator’ is 0), the ‘words’ count will also be 0. Conversely, when a review includes text, the ‘words’ variable captures the word count, adding a layer of depth to the quantitative assessment.

This cleaned version facilitates a more straightforward approach to understanding the patterns and tendencies in customer feedback without the convolution of qualitative data interpretation.

Table 2: Sample of the Cleaned Dataset

Review ID	Stars	Text Indicator	Words
1	5	0	0
2	5	0	0
3	4	0	0
4	5	0	0
5	4	0	0
6	5	0	0
7	5	1	8
8	5	1	29
9	5	0	0
10	4	0	0

2.3 Data Summary

Two figures have been constructed to offer visual summaries of the cleaned dataset, elucidating patterns in customer reviews at a glance.

2.3.1 Frequency of Text Reviews Across Star Categories

The histogram in Figure 1 categorizes reviews based on the star ratings they correspond to and whether the reviews include text. Reviews with text are indicated in blue, and those without in red. A conspicuous majority of reviews across all star ratings do not include textual feedback, as the prevalence of red bars suggests. This is particularly pronounced in the 5-star category, which also features the highest number of reviews, indicating the store’s generally favorable reception.

Notably, as the star rating decreases, the proportion of text reviews (blue bars) increases relative to those without text. This trend culminates in the 1-star category, where the number of text reviews surpasses that of non-text reviews, suggesting that customers who are dissatisfied are more likely to provide written explanations of their experiences. This observation will

be quantitatively examined through logistic regression modeling to confirm the hypothesized relationship between the star rating and the propensity to leave textual feedback.




Figure 1: Frequency of Text Reviews Varies Across Different Star Categories

2.3.2 Overview of Text Length in Reviews by Star Rating

In Figure 2, we present a box plot that examines the length of text in reviews, filtered by star rating, for those containing written commentary (text indicator=1). The logarithmic scale is employed on the y-axis to manage outliers and maintain readability. This transformation allows us to observe the central tendency and dispersion of text lengths without distortion from extreme values.

The plot suggests a tendency for reviews with higher star ratings to have shorter text lengths, although the trend does not exhibit a stark contrast across the star categories. The presence of outliers is apparent, particularly in the lower star categories. To further investigate and model the relationship between the star rating and the length of text, a multilevel regression analysis will be employed, taking into account the variability within and between the different levels of star ratings.

The synthesis of these visuals into our analysis provides preliminary insights into the nature of the reviews. They serve to highlight the initial observations that customers tend to write longer texts when they are less satisfied and tend to leave reviews without text when more




Figure 2: Text Length Varies Across Text Reviews from Different Star Categories

satisfied, with the possible implications of these trends to be explored in subsequent modeling efforts.

2.4 Measurement

The pathway from individual dining experiences to structured dataset entries commences with patrons articulating their encounters at Pizza Hut’s Union Place branch in Colombo via online review platforms. These digital expressions, manifested through both ratings and textual reviews, underpin the dataset utilized in this analysis.

The genesis of this dataset was adeptly orchestrated by its author through the utilization of web scraping tools, specifically using Apify’s Tool for Google Reviews. This approach ensured an aggregation of customer feedback specific to this Pizza Hut location. In an effort to uphold ethical standards and privacy, the dataset author performed preliminary cleaning to anonymize the data by removing personal identifiers from the reviews. This step was critical in safeguarding individual privacy while retaining the dataset’s analytical value.

The kaggle website hosting this dataset offers limited information on its compilation process. However, through direct communication with the dataset author, I gained insight into how the dataset was developed and subsequently created a datasheet to document these methodologies.

My contribution to the dataset’s preparation involved additional cleaning to align it with the study’s analytical goals. This included the introduction of a binary variable (`text_indicator`) to denote the presence of textual feedback and calculating the word count (`words`) for each review. These modifications allowed for a detailed examination of the interplay between customer satisfaction, as evidenced by star ratings, and the likelihood and elaborateness of textual feedback.

3 Models

Model Diagnostics are included in [Appendix A](#).

3.1 Logistic Regression Model for Text Indicator and Star Ratings

Our analytical endeavor aims to decipher the underlying patterns in customer feedback for a Pizza Hut outlet, specifically examining the propensity of customers to leave textual comments alongside their star ratings. The central thesis of our model is to understand the relationship between the quantitative measure of customer satisfaction (the star ratings) and the qualitative aspect (textual feedback).

The binary nature of our outcome of interest, `text_indicator`, which signifies the presence or absence of text in a review, necessitates a modeling approach that can handle dichotomous data effectively. Logistic regression is particularly suited for this task as it predicts the probability of a binary response based on one or more predictor variables. This model provides the probability that a review contains textual feedback for each level of star rating.

The logistic regression model we are using is designed to examine how star ratings are associated with the presence of textual feedback in customer reviews. Mathematically, the model is expressed as:

$$y_i | \pi_i \sim \text{Bernoulli}(\pi_i) \tag{1}$$

$$\text{logit}(\pi_i) = \beta_0 + \beta_1 \times \text{stars}_i \tag{2}$$

In equation (1) and equation (2):

- y_i is the `text_indicator` variable, which indicates the presence (1) or absence (0) of text in the i -th review.
- π_i is the probability that the i -th review contains text.
- stars_i represents the star rating given by the i -th customer.
- β_0 is the intercept of the model.

- β_1 is the effect size of the star rating on the log-odds of the review containing text.

We set the prior distributions for the intercept β_0 and the coefficient β_1 to be normal with a mean of 0 and a standard deviation of 2.5. This reflects our initial assumption that there is no strong prior belief about the direction or magnitude of the rating’s influence on textual feedback.

Through this logistic regression framework, we aim to quantify the extent to which the star rating influences the probability of textual feedback being provided. This quantification will enable us to discern whether customers who are more satisfied (higher stars) or less satisfied (lower stars) are more inclined to provide detailed textual feedback, thus offering insights into consumer behavior and satisfaction.

By integrating the logistic regression model into our analysis, we expect to unveil patterns that could inform strategies to enhance customer engagement and satisfaction, potentially guiding operational improvements and customer service protocols.

3.2 The Multilevel Model for Text Length and Star Ratings

Our objective extends to examining the detail provided in customer reviews, specifically the length of text in reviews that include textual feedback. To achieve this, we built a multilevel model that allows us to assess how the star ratings, representing different levels of customer satisfaction, relate to the verbosity of the reviews.

We opted for a multilevel model approach due to the inherent group structure within our data: reviews can be naturally divided into groups based on the star rating. This model type accommodates the non-independence of observations within the same group (star rating category) and enables us to capture both the fixed effects of the star ratings and the random effects that account for the variability within each star rating level.

Additionally, the length of the review text, counted as the number of words, typically exhibits overdispersion, where the variance exceeds the mean. The negative binomial distribution is chosen over the simpler Poisson distribution to model such count data with overdispersion. The negative binomial model is particularly suited for this kind of data as it introduces an extra parameter to account for the overdispersion.

The negative binomial model is specified as:

$$y_i | (\mu_i, k) \sim \text{NegBinomial}(\mu_i, k) \quad (3)$$

$$\log(\mu_i) = \beta_0 + \beta_1 \times \text{stars}_i \quad (4)$$

In equation (3) and equation (4):

- y_i represents the count of words in the i -th review's text.
- μ_i is the expected count of words for the i -th review.
- k is the dispersion parameter of the negative binomial distribution.
- β_0 and β_1 are the model coefficients for the intercept and the star ratings, respectively.

The multilevel negative binomial model allows for random effects across different levels of star ratings, which can account for variability within each star category:

$$y_i | (\mu_i, k) \sim \text{NegBinomial}(\mu_i, k) \quad (5)$$

$$\log(\mu_i) = \beta_0 + u_{0[\text{stars}_i]} \quad (6)$$

$$u_{0j} \sim \text{Normal}(0, \sigma_{u_0}^2) \quad (7)$$

In equation (5), equation (6), and equation (7):

- $u_{0[\text{stars}_i]}$ is the random intercept for the star rating category to which the i -th review belongs.
- $\sigma_{u_0}^2$ is the variance of the random intercepts across different star rating categories.

Each model was fit using `stan_glm` and `stan_glmr` functions from the `rstanarm` package in R. The models' outputs were then saved as RDS files for further analysis.

By focusing on the subset of data where text indicator is 1, we concentrate on those reviews that actually contain text. This filtering ensures that the model is trained on the most informative data for our research question—understanding the depth of feedback in relation to customer satisfaction levels.

Through this multilevel negative binomial model, we seek to elucidate the patterns that may exist between the satisfaction indicated by the star rating and the length of textual feedback provided. The analysis will explore whether customers who report different levels of satisfaction tend to leave more or less detailed reviews, as reflected in the word count.

By adopting this modeling approach, we aim to extract substantive insights from the text reviews, which are particularly valuable to the store for understanding customer experiences in depth and using those insights for service improvement.

4 Results

4.1 The Relationship Between Text reviews and Star Ratings

The results of our logistic regression model are shown in Table 3 and Figure 3, which together help explain the relationship between the star ratings given by customers and their likelihood of leaving a text review. The relationship between.

Table 3 shows the estimated coefficients of the logistic regression model. The intercept as a baseline comparison shows a positive coefficient, although this coefficient is not statistically significant given the accompanying standard errors. The coefficients for stars=2 through stars=5 are all negative, with factors stars=3, factor stars=4, and factor stars=5 showing more significant coefficients (indicated by smaller standard errors).

The negative coefficient on star rating indicates that customers who give higher star ratings (2 to 5) are less likely to leave text reviews compared to the baseline (implied to be 1 star because not shown). The size of these coefficients increases with star rating, indicating that the likelihood of leaving a text review decreases with increasing satisfaction (star rating).

Figure 3 shows the 90% confidence intervals for the model coefficients. These intervals provide a range of values within which we can be 90% confident about the true value of the coefficient. If the interval of the coefficient does not cross the zero line, it indicates a significant effect at the 10% significance level.

The confidence intervals for stars=3, stars=4, and stars=5 are well below zero, which reinforces the conclusion in Figure 3 that as the star rating increases, the likelihood of a text review is lower. The intervals for stars=2 overlap to zero, indicating less certainty in the effect of a 2-star rating compared to the baseline.

Taken together, these results support the assertion that customers who provide lower star ratings are more likely to leave text reviews, while customers who provide higher star ratings are more likely to leave reviews without text. This pattern can be observed by increasingly negative coefficients for higher stars and their corresponding confidence intervals.

The R-squared values in Table 3 indicate that the proportion of variance in the response variable explained by the model is very low, suggesting that factors other than star ratings may influence the decision to leave a text review.

In summary, logistic regression analysis provides evidence to support the notion that customer satisfaction, as measured by star ratings, is inversely related to the probability of leaving text feedback. This insight is valuable for understanding customer engagement and can inform strategies to encourage more detailed feedback from highly satisfied customers.

Table 3: Whether Customers are Likely to Give a Text Review Based on the Star Rating they Gave

	Give a Text Review
(Intercept)	0.197 (0.166)
factor(stars)2	−0.463 (0.271)
factor(stars)3	−0.872 (0.204)
factor(stars)4	−0.911 (0.180)
factor(stars)5	−1.010 (0.170)
Num.Obs.	4000
R2	0.011
Log.Lik.	−2508.704
ELPD	−2513.8
ELPD s.e.	22.2
LOOIC	5027.5
LOOIC s.e.	44.5
WAIC	5027.5
RMSE	0.47




Figure 3: Credible Intervals for Predictors of Giving a Text Review

4.2 The Different Length of Text Reviews Between Different Stars

The multilevel negative binomial model’s findings are encapsulated in Table 4 and Figure 4. These visualizations allow us to interpret the relationship between the star ratings and the length of the text reviews provided by customers.

Table 4 contrasts the coefficients from the negative binomial model with those from the multilevel negative binomial model. Notably, all coefficients for stars=2 through stars=5 are negative, indicating that as star ratings increase, the expected count of words in reviews tends to decrease. The coefficients are statistically significant, as suggested by the small standard errors, particularly for ratings of 3 stars and above in both models. This aligns with the hypothesis that lower star ratings are associated with longer, presumably more detailed, text reviews.

Figure 4 offers a visual representation of the distribution of the length of text reviews across different star ratings. The plots suggest a clear trend: the group with a 1-star rating shows a distribution that extends towards higher word counts, indicative of longer reviews. As the star ratings increase, the distributions shift leftward, denoting shorter reviews. This pattern is consistent across all star levels and corroborates the conclusion drawn from the coefficients in Table 4.

The combined evidence from the coefficients and the distribution plots substantiates the con-

Table 4: Whether a Customer is Likely to Give a Text Review Based on the Star Rating He Gave

	Neg binomial	Multilevel Neg Binomial
(Intercept)	3.647 (0.116)	2.854 (0.260)
factor(stars)2	-0.413 (0.205)	
factor(stars)3	-0.909 (0.161)	
factor(stars)4	-1.042 (0.132)	
factor(stars)5	-1.335 (0.125)	
Num.Obs.	1308	1308
Log.Lik.	-4669.934	-4669.796
ELPD	-4680.4	-4679.8
ELPD s.e.	55.6	55.4
LOOIC	9360.8	9359.6
LOOIC s.e.	111.3	110.7
WAIC	9360.7	9360.4
RMSE	22.73	22.73

clusion that customers leaving lower star ratings tend to provide more extended textual feedback in their reviews, while those with higher star ratings are inclined to leave briefer comments. This trend may reflect a tendency for less satisfied customers to elaborate on the reasons for their dissatisfaction.

The consistency between the coefficients' significance and the visual trends in review lengths across star ratings enhances the credibility of this conclusion. The models collectively indicate a meaningful and inverse relationship between satisfaction levels, as expressed by star ratings, and the extent of customer engagement in providing textual feedback.

5 Discussion

5.1 Overview of Works

In this article, we conduct a preliminary analysis of customer feedback at Pizza Hut stores in Sri Lanka, focusing on the relationship between star ratings and customers' propensity to leave text reviews and the length of these reviews. Using Kaggle's dataset, we carefully cleaned

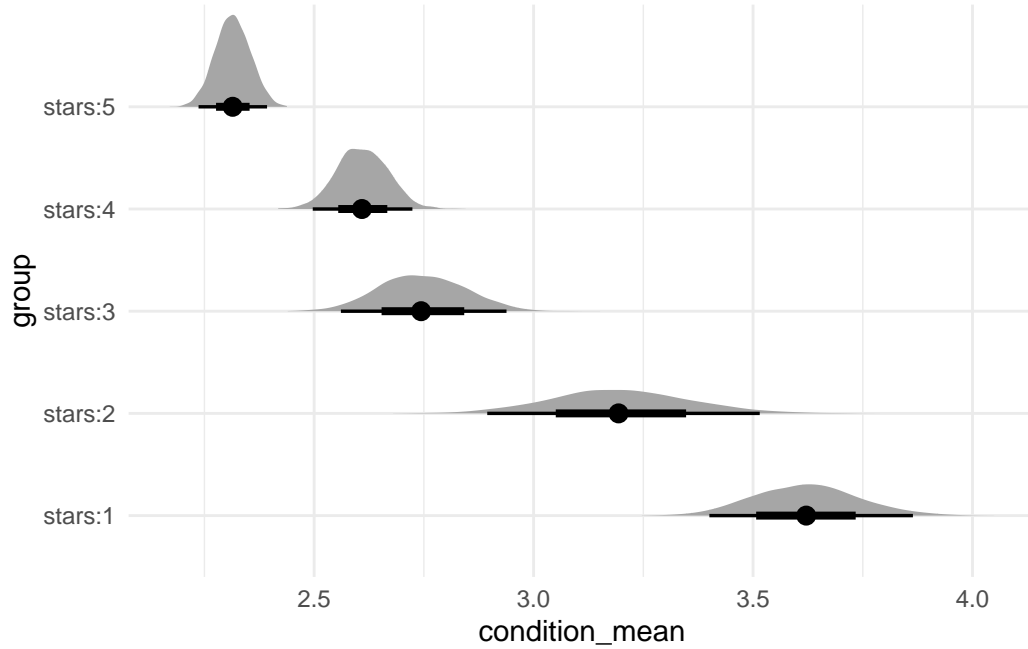


Figure 4: Examining the Distribution of Length of Text for Different Stars

and prepared the data to build two different statistical models. The first model is a logistic regression designed to evaluate the likelihood of a customer leaving text feedback along with the star rating they provide. The second model is a multilevel negative binomial model used to understand the relationship between star ratings and text length in reviews containing written feedback. By analyzing the data through these models, we found that customers who gave lower star ratings were more likely to give reviews with text, and the texts they gave were longer and more detailed than those of customers who gave high star ratings. From this, we seek to uncover the nuances of customer satisfaction and engagement with the aim of providing actionable insights that enhance customer experience and business practices.

5.2 Insights into Societal Behaviors

This paper offers insights into how people behave when it comes to customer service feedback. It uncovers a consistent pattern where customers who aren't happy are more likely to leave detailed text feedback. This seems to be a way for them to express and justify their dissatisfaction. On the other hand, customers who are satisfied usually just leave a rating without writing much. This shows a common behavior where people feel a stronger need to talk about bad experiences, hoping to get recognized and to see changes made. Good experiences, though, tend to be briefly mentioned. This observation emphasizes how it works for businesses to pay

attention to negative feedback since it often holds key insights for improving services and making customers happier.

The study also opens up a conversation about the less obvious aspects of customer satisfaction. While customers who are unhappy tend to explain their issues in detail, those who are content might not say much beyond giving a high rating. This pattern suggests that businesses might miss out on understanding how many customers are actually satisfied, possibly leading to a mistaken view of the overall quality of service. It's a gentle reminder that feedback isn't always loud, and that positive feedback might not always be as visible in customer reviews.

5.3 Reflecting on the Limitations

One of the limitations in our approach is the assumption that the length of the text reviews is directly proportional to the level of detail and richness of the content. However, a longer review does not always equate to more substantive feedback. Some lengthy reviews may be verbose without offering concrete insights, while shorter reviews could be concise yet rich in valuable information. This paper's methodology does not account for the semantic depth or the specific qualitative insights that might be present in the text, which requires a more nuanced analysis, potentially through natural language processing techniques.

Additionally, our analysis is constrained to a single Pizza Hut branch, which limits the generalizability of the findings across different locations, cultures, or even other service industries. Customer behavior can be highly context-specific, and what holds for one branch in Sri Lanka may not be true elsewhere. Also, the scope of the data does not capture the non-linear complexities of customer experiences—factors such as customer expectations, individual differences, and specific circumstances surrounding each visit are beyond the reach of this study. These elements represent areas for further research to build on the preliminary findings presented here.

5.4 Future Directions and Unexplored Avenues

The path forward from this study points to the integration of natural language processing (NLP) to unravel the qualitative depth of customer reviews. Future research could deploy sentiment analysis to gauge the emotional tone and extract specific themes from textual feedback, providing a more granular understanding of customer satisfaction. Expanding the dataset to encompass multiple locations and a broader range of service industries would also offer a more comprehensive view of customer feedback behaviors across different contexts. Additionally, investigating the impact of individual customer characteristics, such as loyalty and prior experiences, could yield insights into the nuanced drivers of review behaviors. Pursuing these avenues will enhance our understanding of customer feedback mechanisms and inform the development of more targeted strategies for improving customer experience and engagement.

Appendix

A Model Diagnostics

A.1 Diagnostics for the Logistic Regression Model

Figure 5a is a trace plot for the logistic regression model discussed in Section 3.1. It illustrates the sampled values across multiple iterations of the Markov Chain Monte Carlo (MCMC) process for each model parameter. The plot shows that the chains for each parameter are well-mixed and stable, without noticeable drifts or divergences, suggesting good convergence. This indicates that the sampling process and the estimates for the model parameters are reliable.

Figure 5b is an Rhat plot for the logistic regression model in Section 3.1. It displays the Rhat statistic for each chain, which is a measure of convergence. An Rhat value of 1 indicates perfect convergence, while values greater than 1 suggest that more sampling might be necessary. The plot shows that all Rhat values are close to 1, none exceeding the threshold of 1.1, which suggests that convergence has been achieved, and the chains have sampled from the posterior distribution effectively.




Figure 5: Checking the Convergence of the MCMC Algorithm for the Logistic Regression Model

A.2 Diagnostics for the Multilevel Model

Figure 6a is a trace plot for the multilevel model referenced in Section 3.2. This graph shows the behavior of the sampled values across iterations for the model's parameters. We observe that while most parameters have traces that suggest convergence, with dense and intermingled lines, there are a few, notably the 'reciprocal_dispersion' parameter, that exhibit less stable behavior, with wider fluctuations. This instability suggests potential issues with convergence

for some parameters, indicating that the model may require further tuning or additional iterations to ensure reliable parameter estimates.

Figure 6b is an Rhat plot for the multilevel model in Section 3.2. It displays Rhat values for the model’s parameters, with values ideally close to 1.00 indicating satisfactory convergence. However, the plot indicates that one of the parameter’s Rhat values exceeds the acceptable threshold, represented by a value greater than 1.10. This is a clear indication that the model has not converged well for this particular parameter, signaling the need for additional diagnostic checks or extended sampling to achieve convergence.

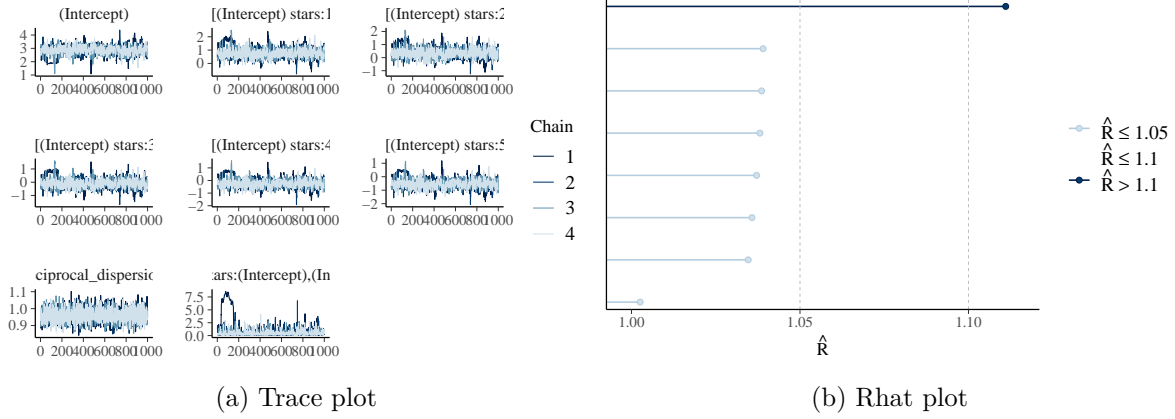


Figure 6: Checking the Convergence of the MCMC Algorithm for the Multilevel Model

References

- Arel-Bundock, Vincent. 2022. *Modelsummary: Summary Tables and Plots for Statistical Models and Data*. <https://CRAN.R-project.org/package=modelsummary>.
- Bolker, Ben et al. 2022. *Broom.mixed: Tidying Methods for Mixed Models*. <https://CRAN.R-project.org/package=broom.mixed>.
- Foundation, Apache Software. 2022. *Arrow: Integration to 'Apache Arrow'*. <https://CRAN.R-project.org/package=arrow>.
- Kaggle. 2022. “Kaggle API.” <https://www.kaggle.com/docs/api>.
- Kay, Matthew. 2022a. *Ggdist: Visualizations of Distributions and Uncertainty*. <https://mjskay.github.io/ggdist/>.
- . 2022b. *Tidybayes: Tidy Data and Geoms for Bayesian Models*. <https://mjskay.github.io/tidybayes/>.
- Leeper, Thomas. 2022. *Dataverse: Client for Dataverse 4 Repositories*. <https://CRAN.R-project.org/package=dataverse>.
- Neuwirth, Erich. 2022. *RColorBrewer: ColorBrewer Palettes*. <https://CRAN.R-project.org/package=RColorBrewer>.

- R Core Team. 2022. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Team, Stan Development. 2022. *Rstanarm: Bayesian Applied Regression Modeling via Stan*. <https://mc-stan.org/rstanarm>.
- Wickham, Hadley. 2022. *Ggplot2: Create Elegant Data Visualisations Using the Grammar of Graphics*. <https://CRAN.R-project.org/package=ggplot2>.
- Wickham, Hadley, and et al. 2019. *Tidyverse: Easily Install and Load the Tidyverse*. <https://CRAN.R-project.org/package=tidyverse>.
- . 2022. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.
- Xie, Yihui. 2014. *Knitr: A General-Purpose Package for Dynamic Report Generation in r*. <https://yihui.org/knitr/>.